

Part (B)

The `check_grad` function is used and the following is the output.

Gradient check using `check_grad`: 6.093988195535878e-06

Gradient check using `approx_fprime`:

Value 0: -0.014803935965964259

Value 1: -0.014803935965964259

Value 2: -0.014978769669034562

Value 3: -0.01480393596598994

Value 4: -0.01480393596598994

.

.

.

Value 59205: 4.5192294351136574e-05

Value 59206: -0.13773177587326604

Value 59207: 0.00014157508498854605

Value 59208: 8.768985537933954e-06

Value 59209: -0.00013831824574581144

```
print(scipy.optimize.check_grad(
    lambda weights: fCE(trainX, trainY, weights, NUM_HIDDEN_LAYERS, NUM_INPUT, NUM_HIDDEN, NUM_OUTPUT), # No slicing on trainX
    lambda weights: gradCE(trainX, trainY, weights, NUM_HIDDEN_LAYERS, NUM_INPUT, NUM_HIDDEN, NUM_OUTPUT),
    weights
))

weights, avg_loss = train(X_train, y_train, weights, X_val, y_val, lr, num_epochs, batch_size, alpha, NUM_HIDDEN_LAYERS, NUM_INPUT, NUM_HIDDEN, NUM_OUTPUT) #
val_accuracy = validate(X_val, y_val, weights, NUM_HIDDEN_LAYERS, NUM_INPUT, NUM_HIDDEN, NUM_OUTPUT)
test_accuracy = test(X_test, y_test, weights, NUM_HIDDEN_LAYERS, NUM_INPUT, NUM_HIDDEN, NUM_OUTPUT)

print(scipy.optimize.approx_fprime(weights, lambda weights: fCE(trainX, trainY, weights, NUM_HIDDEN_LAYERS, NUM_INPUT, NUM_HIDDEN, NUM_OUTPUT), 1e-6))
```

✓ 46.9s Python

3.5305727448724214e-06
[-1.48039360e-02 -1.48039360e-02 -1.49787697e-02 ... 1.41575085e-04
8.76898554e-06 -1.38318246e-04]

All the values are saved in **Part_B_grad_weight.txt**

Part (C)

```
training_data = np.reshape(load("fashion-mnist/train-shapes.npy"), (-1, 28*28))
training_labels = np.load("fashion-mnist/train-labels.npy")

num_datapoints = training_data.shape[0]
split_index = int(0.9*num_datapoints)
indices = np.arange(num_datapoints)
np.random.shuffle(indices)

train_indices = indices[:split_index]
val_indices = indices[split_index:]

X_train, X_val = training_data[train_indices], training_data[val_indices]
y_train, y_val = training_labels[train_indices], training_labels[val_indices]

X_train = X_train / 255.0 - 0.5
X_val = X_val / 255.0 - 0.5

best_hyp, best_weights, best_acc = validation(X_train, y_train, X_val, y_val)
```

Python

Hidden Lrs: 3, n_hidden_units: 30, Num Epoch: 50, Batch size: 16, Learning rate: 0.005, Unregularise CE: 0.32030, Test accuracy: 88.42%
Hidden Lrs: 3, n_hidden_units: 30, Num Epoch: 100, Batch size: 16, Learning rate: 0.005, Unregularise CE: 0.20803, Test accuracy: 88.33%
Hidden Lrs: 3, n_hidden_units: 30, Num Epoch: 50, Batch size: 32, Learning rate: 0.005, Unregularise CE: 0.20984, Test accuracy: 88.08%
Hidden Lrs: 3, n_hidden_units: 30, Num Epoch: 100, Batch size: 32, Learning rate: 0.005, Unregularise CE: 0.31040, Test accuracy: 88.31%
Hidden Lrs: 3, n_hidden_units: 30, Num Epoch: 50, Batch size: 64, Learning rate: 0.005, Unregularise CE: 0.30692, Test accuracy: 87.18%
Hidden Lrs: 3, n_hidden_units: 30, Num Epoch: 100, Batch size: 64, Learning rate: 0.005, Unregularise CE: 0.48018, Test accuracy: 87.94%
Hidden Lrs: 3, n_hidden_units: 30, Num Epoch: 50, Batch size: 16, Learning rate: 0.01, Unregularise CE: 0.28526, Test accuracy: 87.68%
Hidden Lrs: 3, n_hidden_units: 30, Num Epoch: 100, Batch size: 16, Learning rate: 0.01, Unregularise CE: 0.22049, Test accuracy: 88.13%
Hidden Lrs: 3, n_hidden_units: 30, Num Epoch: 50, Batch size: 32, Learning rate: 0.01, Unregularise CE: 0.32780, Test accuracy: 88.27%
Hidden Lrs: 3, n_hidden_units: 30, Num Epoch: 100, Batch size: 32, Learning rate: 0.01, Unregularise CE: 0.26238, Test accuracy: 87.74%
Hidden Lrs: 3, n_hidden_units: 30, Num Epoch: 50, Batch size: 64, Learning rate: 0.01, Unregularise CE: 0.48073, Test accuracy: 87.92%
Hidden Lrs: 3, n_hidden_units: 30, Num Epoch: 100, Batch size: 64, Learning rate: 0.01, Unregularise CE: 0.32283, Test accuracy: 88.32%
Hidden Lrs: 3, n_hidden_units: 30, Num Epoch: 50, Batch size: 16, Learning rate: 0.05, Unregularise CE: 0.24988, Test accuracy: 88.12%
Hidden Lrs: 3, n_hidden_units: 30, Num Epoch: 100, Batch size: 16, Learning rate: 0.05, Unregularise CE: 0.20744, Test accuracy: 87.59%
Hidden Lrs: 3, n_hidden_units: 30, Num Epoch: 50, Batch size: 32, Learning rate: 0.05, Unregularise CE: 0.25572, Test accuracy: 87.88%
Hidden Lrs: 3, n_hidden_units: 30, Num Epoch: 100, Batch size: 32, Learning rate: 0.05, Unregularise CE: 0.20340, Test accuracy: 87.47%
Hidden Lrs: 3, n_hidden_units: 30, Num Epoch: 50, Batch size: 64, Learning rate: 0.05, Unregularise CE: 0.28187, Test accuracy: 87.47%
Hidden Lrs: 3, n_hidden_units: 30, Num Epoch: 100, Batch size: 64, Learning rate: 0.05, Unregularise CE: 0.22720, Test accuracy: 88.30%
Hidden Lrs: 3, n_hidden_units: 40, Num Epoch: 50, Batch size: 16, Learning rate: 0.005, Unregularise CE: 0.32188, Test accuracy: 87.61%
Hidden Lrs: 3, n_hidden_units: 40, Num Epoch: 100, Batch size: 16, Learning rate: 0.005, Unregularise CE: 0.22273, Test accuracy: 87.90%
Hidden Lrs: 3, n_hidden_units: 40, Num Epoch: 50, Batch size: 32, Learning rate: 0.005, Unregularise CE: 0.30352, Test accuracy: 87.69%
Hidden Lrs: 3, n_hidden_units: 40, Num Epoch: 100, Batch size: 32, Learning rate: 0.005, Unregularise CE: 0.33339, Test accuracy: 88.41%
Hidden Lrs: 3, n_hidden_units: 40, Num Epoch: 50, Batch size: 64, Learning rate: 0.005, Unregularise CE: 0.30202, Test accuracy: 88.64%
Hidden Lrs: 3, n_hidden_units: 40, Num Epoch: 100, Batch size: 64, Learning rate: 0.005, Unregularise CE: 0.39749, Test accuracy: 88.47%
Hidden Lrs: 3, n_hidden_units: 40, Num Epoch: 50, Batch size: 16, Learning rate: 0.01, Unregularise CE: 0.27828, Test accuracy: 88.17%
Hidden Lrs: 3, n_hidden_units: 50, Num Epoch: 50, Batch size: 32, Learning rate: 0.05, Unregularise CE: 0.28024, Test accuracy: 88.53%
Hidden Lrs: 3, n_hidden_units: 50, Num Epoch: 100, Batch size: 32, Learning rate: 0.05, Unregularise CE: 0.17256, Test accuracy: 87.77%
Hidden Lrs: 3, n_hidden_units: 50, Num Epoch: 50, Batch size: 64, Learning rate: 0.05, Unregularise CE: 0.26747, Test accuracy: 88.13%
Hidden Lrs: 3, n_hidden_units: 50, Num Epoch: 100, Batch size: 64, Learning rate: 0.05, Unregularise CE: 0.13745, Test accuracy: 88.88%

Output is truncated. View as a [scrollable element](#) or open in a [text editor](#). Adjust cell output settings...

```
np.save("best_model_weights_hp.npy", best_weights)
np.save("best_model_hyperparameters_hp.npy", best_hyp)
np.save("best_model_accuracy_hp.npy", best_acc)
```

Python

best_acc=100


Python

88.94166666666666

Accuracy: 88.94166666666666

Accuracy for each iteration is saved in **hyperparameter_tune.txt**

Part (D)



```
In [48]: # Import the NumPy
import numpy as np

# Import the 'optimize' module from SciPy
import scipy.optimize

# Importing plt from matplotlib.pyplot
import matplotlib.pyplot as plt
```

Unpack Function

Function to unpack a single vector of weights into weight matrices and bias vectors for a multi-layer neural network.

Arguments:

- weights: Flattened array containing all weights and biases.
- NUM_HIDDEN_LAYERS: Number of hidden layers in the neural network.
- NUM_INPUT: Number of input features.
- NUM_HIDDEN: List specifying the number of units in each hidden layer.
- NUM_OUTPUT: Number of output units.

Returns:

- Ws: List of weight matrices for each layer.
- bs: List of bias vectors for each layer.

```
In [17]: def unpack(weights, NUM_HIDDEN_LAYERS, NUM_INPUT, NUM_HIDDEN, NUM_OUTPUT):

    # Initialize list to store weight matrices
    Ws = []

    # Unpack the weight matrix for the input layer to the first hidden layer
    start = 0
    end = NUM_INPUT * NUM_HIDDEN[0] # Calculate the range for the first wei
    W = weights[start:end] # Slice the weight vector for the first layer's
    Ws.append(W) # Append the weight vector to the list

    # Unpack weight matrices between hidden layers
    for i in range(NUM_HIDDEN_LAYERS - 1): # Loop through hidden layers
        start = end # Update the start index
        end = end + NUM_HIDDEN[i] * NUM_HIDDEN[i+1] # Calculate the end ind
        W = weights[start:end] # Slice the weight vector for the current la
        Ws.append(W) # Append the weight vector to the list

    # Unpack the weight matrix for the last hidden layer to the output layer
    start = end # Update the start index for the last layer
    end = end + NUM_HIDDEN[-1] * NUM_OUTPUT # Calculate the end index for t
    W = weights[start:end] # Slice the weight vector for the output layer's
    Ws.append(W) # Append the weight vector to the list
```

```

# Reshape the weight vectors into matrices for each layer
Ws[0] = Ws[0].reshape(NUM_HIDDEN[0], NUM_INPUT) # Reshape the first layer
for i in range(1, NUM_HIDDEN_LAYERS): # Loop through hidden layers
    Ws[i] = Ws[i].reshape(NUM_HIDDEN[i], NUM_HIDDEN[i-1]) # Reshape the
Ws[-1] = Ws[-1].reshape(NUM_OUTPUT, NUM_HIDDEN[-1]) # Reshape the final

# Initialize list to store bias vectors
bs = []

# Unpack bias vector for the first hidden layer
start = end # Set start index for biases
end = end + NUM_HIDDEN[0] # Calculate the end index for the first bias
b = weights[start:end] # Slice the weight vector for the biases
bs.append(b) # Append the bias vector to the list

# Unpack bias vectors for the hidden layers
for i in range(NUM_HIDDEN_LAYERS - 1): # Loop through hidden layers
    start = end # Update the start index for the next bias vector
    end = end + NUM_HIDDEN[i+1] # Calculate the end index for the next
    b = weights[start:end] # Slice the weight vector for the biases
    bs.append(b) # Append the bias vector to the list

# Unpack the bias vector for the output layer
start = end # Set start index for the output layer's bias
end = end + NUM_OUTPUT # Calculate the end index for the output layer's
b = weights[start:end] # Slice the weight vector for the biases
bs.append(b) # Append the bias vector to the list

# Return the unpacked weight matrices and bias vectors
return Ws, bs

```

Initializing the weights and biases

Initializes the weights and biases for a multi-layer neural network using a modified Kaiming He Uniform initialization for weights and a small constant for biases.

Arguments:

- NUM_HIDDEN_LAYERS: Number of hidden layers in the neural network.
- NUM_INPUT: Number of input features.
- NUM_HIDDEN: List specifying the number of units in each hidden layer.
- NUM_OUTPUT: Number of output units.

Returns:

- Ws: List of initialized weight matrices.
- bs: List of initialized bias vectors.

```
In [18]: def initWeightsAndBiases(NUM_HIDDEN_LAYERS, NUM_INPUT, NUM_HIDDEN, NUM_OUTPUT)
```

```

Ws = [] # List to store weight matrices
bs = [] # List to store bias vectors

# Strategy:
# - Weight initialization: Kaiming He Uniform initialization, suitable for
# - Bias initialization: Small positive constant (0.01).

# Set random seed for reproducibility
np.random.seed(0)

# Initialize weights between input layer and the first hidden layer
W = 2 * (np.random.random(size=(NUM_HIDDEN[0], NUM_INPUT)).astype(np.float32))
Ws.append(W) # Add the initialized weight matrix to the list

# Initialize biases for the first hidden layer
b = 0.01 * np.ones(NUM_HIDDEN[0], dtype=np.float32) # Bias initialized
bs.append(b) # Add the initialized bias vector to the list

# Initialize weights and biases for hidden layers
for i in range(NUM_HIDDEN_LAYERS - 1):
    # Initialize weights between the ith hidden layer and the next hidden layer
    W = 2 * (np.random.random(size=(NUM_HIDDEN[i+1], NUM_HIDDEN[i])).astype(np.float32))
    Ws.append(W) # Add the weight matrix to the list

    # Initialize biases for the next hidden layer
    b = 0.01 * np.ones(NUM_HIDDEN[i+1], dtype=np.float32) # Bias initialized
    bs.append(b) # Add the bias vector to the list

# Initialize weights for the final output layer (from last hidden layer)
W = 2 * (np.random.random(size=(NUM_OUTPUT, NUM_HIDDEN[-1])).astype(np.float32))
Ws.append(W) # Add the weight matrix to the list

# Initialize biases for the output layer
b = 0.01 * np.ones(NUM_OUTPUT, dtype=np.float32) # Bias initialized
bs.append(b) # Add the bias vector to the list

# Return the lists of weight matrices and bias vectors
return Ws, bs

```

In [27]: # Helper functions

```

def relu(z):
    return np.maximum(0, z)

def relu_derivative(z):
    return (z > 0).astype(float)

def softmax(Z):
    # to avoid large exponent values, we will use the deviation idealogy of
    exponent_z = np.exp(Z - np.max(Z,axis=1,keepdims=True))

    prediction = exponent_z/(np.sum(exponent_z,axis=1,keepdims=True))

    return prediction

```

```

def fCE (X, Y, weights):
    Ws, bs = unpack(weights)
    H_activation, _ = forward_pass(X, Ws, bs, NUM_HIDDEN_LAYERS)
    y_pred = H_activation[-1]
    ce = loss(Y, y_pred, Ws, 0)
    return ce

def gradCE(X, Y, weights):
    # Unpack the flattened weights and biases
    Ws, bs = unpack(weights)

    wt_grad = [] # List to store weight gradients
    bi_grad = [] # List to store bias gradients

    # Perform a forward pass through the network
    H_activation, z_score = forward_pass(X, Ws, bs, NUM_HIDDEN_LAYERS)

    # Initialize one-hot encoded labels (y_onehot)
    y_onehot = np.zeros_like(H_activation[-1])

    # Convert Y to integer type if needed
    Y = Y.astype(int)

    # One-hot encode the ground truth labels
    y_onehot[np.arange(len(Y)), Y] = 1

    # Compute the gradient of the loss w.r.t. the output layer (softmax layer)
    g = H_activation[-1] - y_onehot # Difference between prediction and ground truth

    # Backpropagate through each layer, starting from the output layer
    for i in range(NUM_HIDDEN_LAYERS, -1, -1):
        # Gradient of weights ( $g^T * H$ ) and biases (sum of gradients over batch)
        grad_wt = np.dot(g.T, H_activation[i]) / len(Y)
        grad_bi = np.sum(g, axis=0, keepdims=True) / len(Y)

        # Append the computed gradients to the respective lists
        wt_grad.append(grad_wt)
        bi_grad.append(grad_bi)

        # If not the first layer (input layer)
        # Backpropagate the gradient to the previous layer using the ReLU derivative
        if i > 0:
            g = np.dot(g, Ws[i]) * relu_derivative(z_score[i-1])

    # Reverse the gradients list so they align correctly with the weight matrices
    wt_grad.reverse()
    bi_grad.reverse()

    # Flatten and concatenate all gradients (weights and biases) into a single vector
    allGradientsAsVector = np.hstack([W.flatten() for W in wt_grad] + [b.flatten() for b in bi_grad])

    return allGradientsAsVector

```

Loss Function

Computes the cross-entropy loss between the true labels (`y_label`) and the predicted probabilities (`y_pred`) for a batch of data. Optionally, it also supports L2 regularization.

Arguments:

- `y_label`: Ground truth labels (as integer indices).
- `y_pred`: Predicted probabilities from the softmax function.
- `W`: Weight matrix (used for optional L2 regularization, currently not used).
- `alpha`: Regularization parameter (L2 regularization strength, currently not used).

Returns:

- `loss`: The average cross-entropy loss for the batch.

```
In [28]: def loss(y_label, y_pred, W, alpha):  
  
    # Get the batch size, i.e., number of samples in the current batch.  
    batch_s = y_label.shape[0] # This ensures loss is computed as an average  
  
    # Ensure that y_label is of integer type, as it represents class indices  
    y_label = y_label.astype(int)  
  
    # Compute the negative log-likelihood (cross-entropy loss) for the true  
    # For each sample, we access the predicted probability for the true label  
    prob = -np.log(y_pred[range(batch_s), y_label]) # Select the predicted  
  
    # Compute the total loss by summing all the individual losses, and then  
    loss = np.sum(prob) / batch_s # Average loss across the batch.  
  
    # Optionally, L2 regularization (currently commented out).  
    # It penalizes large weights to prevent overfitting, and helps to smooth  
    # reg_loss = alpha / 2 * np.sum(np.square(W)) # L2 regularization term  
  
    # Compute total batch loss (currently not using regularization).  
    # batch_loss = loss + reg_loss # Total loss (including regularization).  
  
    return loss # Return the average cross-entropy loss (without regularization)
```

Forward Pass

Performs the forward pass through a neural network, computing activations for each layer.

Arguments:

- `x`: Input data (batch of samples).
- `Weights`: List of weight matrices for each layer.
- `bias`: List of bias vectors for each layer.

- NUM_HIDDEN_LAYERS: Number of hidden layers in the network.

Returns:

- H_activation: A list containing the activations for each layer (including input and output layers).
- z_preactivation_score: A list containing the pre-activation scores (z values) for each layer.

```
In [29]: def forward_pass(x, Weights, bias, NUM_HIDDEN_LAYERS):

    z_preactivation_score = [] # To store the pre-activation (z) values for
    H_activation = [x] # List to store activations, starting with the input

    # Loop over each hidden layer to compute pre-activations (z) and activations
    for i in range(NUM_HIDDEN_LAYERS):
        # Compute the pre-activation score for the i-th layer: z = H_prev *
        z = np.dot(H_activation[-1], Weights[i].T) + bias[i]

        # Store the pre-activation score (before applying the activation function)
        z_preactivation_score.append(z)

        # Apply the ReLU activation function to the pre-activation score to get
        H = relu(z)

        # Append the activation (H) to the list of activations.
        H_activation.append(H)

    # For the output layer (final layer), apply softmax to get predicted probabilities
    z = np.dot(H_activation[-1], Weights[-1].T) + bias[-1] # Pre-activation score

    # Store the pre-activation score for the output layer.
    z_preactivation_score.append(z)

    # Apply softmax to the final pre-activation score to compute the predicted class
    y_pred = softmax(z)

    # Append the output layer activation (y_pred) to the list of activations
    H_activation.append(y_pred)

    # Return both the activations and the pre-activation scores for each layer
    return H_activation, z_preactivation_score
```

Back Propagation

Performs the backpropagation algorithm to compute gradients and update weights and biases using gradient descent with L2 regularization.

Arguments:

- x_batch: Batch of input data.
- y_batch: Batch of true labels (ground truth).

- `h_activations`: List of activations from the forward pass (including input and output).
- `z_score`: List of pre-activation scores (z values) from the forward pass.
- `Weights`: List of weight matrices for each layer.
- `bias`: List of bias vectors for each layer.
- `lr`: Learning rate for gradient descent.
- `alpha`: L2 regularization strength.
- `NUM_HIDDEN_LAYERS`: Number of hidden layers in the neural network.

Returns:

- `Weights`: Updated weight matrices after gradient descent.
- `bias`: Updated bias vectors after gradient descent.

```
In [30]: def backprop(x_batch, y_batch, h_activations, z_score, Weights, bias, lr, alpha):

    wt_gradient = [] # To store gradients for weight matrices
    bias_gradient = [] # To store gradients for bias vectors
    batch_s = y_batch.shape[0] # Get the batch size

    # Create a one-hot encoded matrix for the ground truth labels
    y_onehot = np.zeros_like(h_activations[-1]) # Shape is same as the output
    y_onehot[np.arange(batch_s), y_batch] = 1 # Set the correct label index

    # Compute the initial gradient of the loss with respect to the output (y)
    g = h_activations[-1] - y_onehot # Gradient of cross-entropy loss w.r.t output

    # Loop over each layer, starting from the output layer and going backwards
    for i in range(NUM_HIDDEN_LAYERS, -1, -1):
        # Compute the gradient of the weight matrix for the current layer
        grad_wt = np.dot(g.T, h_activations[i]) / batch_s # Divide by batch size
        grad_bi = np.sum(g, axis=0, keepdims=True) / batch_s # Compute gradient for bias

        # Store the gradients for later use
        wt_gradient.append(grad_wt)
        bias_gradient.append(grad_bi)

        # Propagate the gradient to the previous layer (if not the input layer)
        if i > 0:
            g = np.dot(g, Weights[i]) * relu_derivative(z_score[i - 1]) # Chain rule

    # Reverse the gradients list to maintain the correct order (from first layer to last)
    wt_gradient.reverse()
    bias_gradient.reverse()

    # Gradient descent update step for weights and biases
    for i in range(NUM_HIDDEN_LAYERS + 1):
        # Update weights using gradient descent and L2 regularization (alpha)
        Weights[i] -= lr * (wt_gradient[i] + alpha * Weights[i])
        # Update bias using gradient descent
        bias[i] -= lr * bias_gradient[i].reshape(-1)
```

```
return Weights, bias # Return the updated weights and biases
```

Train Function

Train the neural network using stochastic gradient descent (SGD) with mini-batches.

Arguments:

- X_train: Training data inputs (features).
- y_train: Training data labels (targets).
- weights: Initial flattened weights and biases (single vector).
- X_val: Validation data inputs (optional for validation purposes).
- y_val: Validation data labels (optional for validation purposes).
- lr: Learning rate for gradient descent.
- num_epochs: Number of training epochs.
- batch_size: Size of mini-batches for SGD.
- alpha: L2 regularization strength.
- NUM_HIDDEN_LAYERS: Number of hidden layers in the neural network.
- NUM_INPUT: Number of input features (size of input layer).
- NUM_HIDDEN: List of sizes for each hidden layer.
- NUM_OUTPUT: Number of output classes (size of output layer).

Returns:

- WTS: Flattened final weights and biases after training.
- avg_loss: Average loss over all epochs.

```
In [40]: def train(X_train, y_train, weights, X_val, y_val, lr, num_epochs, batch_size):  
  
    num_tr_samples = X_train.shape[0] # Total number of training samples  
    tr_indices = np.arange(num_tr_samples) # Create an array of sample indices  
  
    # Unpack the initial weights and biases from the flattened vector  
    Weights, bias = unpack(weights, NUM_HIDDEN_LAYERS, NUM_INPUT, NUM_HIDDEN)  
  
    total_loss = 0 # Initialize total loss accumulator  
  
    # Training loop for multiple epochs  
    for epoch in range(num_epochs):  
        np.random.shuffle(tr_indices) # Shuffle the training sample indices  
  
        # Iterate over mini-batches  
        for num in range(0, num_tr_samples, batch_size):  
            # Select a mini-batch of data based on the current indices  
            batch_index = tr_indices[num: num + batch_size]  
            x_batch = X_train[batch_index] # Extract the inputs for the mini-batch  
            y_batch = y_train[batch_index] # Extract the labels for the mini-batch
```

```

        # Perform a forward pass to compute activations and pre-activations
        H_activations, z_list = forward_pass(x_batch, Weights, bias, NUM_HIDDEN_LAYERS)

        # Perform backpropagation and update weights and biases
        Weights, bias = backprop(x_batch, y_batch, H_activations, z_list)

        # After each epoch, calculate the loss for the entire training set
        activations, _ = forward_pass(X_train, Weights, bias, NUM_HIDDEN_LAYERS)
        epoch_loss = loss(y_train, activations[-1], Weights, alpha) # Compute loss

        total_loss += epoch_loss # Accumulate the total loss over epochs

    # After training, flatten the final weights and biases into a single vector
    WTS = np.hstack([W.flatten() for W in Weights] + [b.flatten() for b in biases])

    # Calculate the average loss over all epochs
    avg_loss = total_loss / num_epochs

    return WTS, avg_loss # Return the final weights and average loss

```

Validate Function

Validate the performance of the neural network on a validation dataset.

Arguments:

- X_val: Validation data inputs (features).
- y_val: Validation data labels (targets).
- weights: Current weights and biases of the model (flattened vector).
- NUM_HIDDEN_LAYERS: Number of hidden layers in the neural network.
- NUM_INPUT: Number of input features (size of input layer).
- NUM_HIDDEN: List of sizes for each hidden layer.
- NUM_OUTPUT: Number of output classes (size of output layer).

Returns:

- accuracy: The accuracy of the model on the validation dataset (fraction of correct predictions).

```

In [41]: def validate(X_val, y_val, weights, NUM_HIDDEN_LAYERS, NUM_INPUT, NUM_HIDDEN_LAYERS)

        # Unpack the weights and biases from the flattened weights vector
        Weights, bias = unpack(weights, NUM_HIDDEN_LAYERS, NUM_INPUT, NUM_HIDDEN_LAYERS)

        # Perform forward pass on the validation data to compute activations
        activations, _ = forward_pass(X_val, Weights, bias, NUM_HIDDEN_LAYERS)

        # The last element of activations corresponds to the output layer prediction
        y_pred = activations[-1]

        # Get the predicted class labels by taking the index of the maximum predicted value
        y_pred_labels = np.argmax(y_pred, axis=1)

```

```

# Calculate the accuracy by comparing predicted labels with the true labels
accuracy = np.mean(y_pred_labels == y_val) # Fraction of correct predictions

# Optionally print the validation accuracy (commented out)
# print(f'Validation accuracy: {accuracy * 100:.2f}%')

return accuracy # Return the computed accuracy

```

```

In [42]: def test(X_test, y_test, weights, NUM_HIDDEN_LAYERS, NUM_INPUT, NUM_HIDDEN, NUM_OUTPUT):

    Weights, bias = unpack(weights, NUM_HIDDEN_LAYERS, NUM_INPUT, NUM_HIDDEN, NUM_OUTPUT)
    # Perform forward pass on validation data
    activations, _ = forward_pass(X_test, Weights, bias, NUM_HIDDEN_LAYERS)

    # Assuming the last activation is the output layer
    y_pred = activations[-1]

    y_pred_labels = np.argmax(y_pred, axis=1)

    # Calculate accuracy
    accuracy = np.mean(y_pred_labels == y_test)
    # print(f'Test accuracy: {accuracy * 100:.2f}%')

    return accuracy

```

```

In [44]: if __name__ == "__main__":
    # Load training data.
    # Recommendation: divide the pixels by 255 (so that their range is [0-1])
    # 0.5 (so that the range is [-0.5,+0.5]).

    NUM_HIDDEN_LAYERS = 3
    NUM_INPUT = 784
    NUM_HIDDEN = NUM_HIDDEN_LAYERS * [ 64 ] # you can customise here
    NUM_OUTPUT = 10

    training_data = np.reshape(np.load("fashion_mnist_train_images.npy"), (-1, NUM_INPUT))
    training_labels = np.load("fashion_mnist_train_labels.npy")

    num_datapoints = training_data.shape[0]
    split_index = int(0.8*num_datapoints)
    indices = np.arange(num_datapoints)
    np.random.shuffle(indices)

    train_indices = indices[:split_index]
    val_indices = indices[split_index:]

    X_train, X_val = training_data[train_indices], training_data[val_indices]
    y_train, y_val = training_labels[train_indices], training_labels[val_indices]

    X_train = X_train / 255.0 - 0.5
    X_val = X_val / 255.0 - 0.5

    testing_data = np.reshape(np.load("fashion_mnist_test_images.npy"), (-1, NUM_INPUT))

```

```

testing_labels= np.load("fashion_mnist_test_labels.npy")

X_test = testing_data
y_test = testing_labels
X_test = X_test / 255.0 - 0.5

#     num_epochs = 100
#     batch_size =64
#     alpha =0.0
#     lr = 0.01

num_epochs = 10
batch_size = 8
alpha = 0.0
lr = 0.01

Ws, bs = initWeightsAndBiases(NUM_HIDDEN_LAYERS,NUM_INPUT,NUM_HIDDEN,NUM
weights = np.hstack([ W.flatten() for W in Ws ] + [ b.flatten() for b in
weights = weights.astype(np.float32)

trainX, trainY = X_train[:5].astype(np.float32), y_train[:5].astype(np.f

print(scipy.optimize.check_grad(
    lambda weights_: fCE(trainX, trainY, weights_, NUM_HIDDEN_LAYERS,NUM_
    lambda weights_: gradCE(trainX, trainY, weights_, NUM_HIDDEN_LAYERS,
    weights
))

weights, avg_loss = train(X_train, y_train, weights, X_val, y_val, lr,nu
val_accuracy = validate(X_val, y_val, weights, NUM_HIDDEN_LAYERS,NUM_INF
test_accuracy = test(X_test, y_test, weights, NUM_HIDDEN_LAYERS,NUM_INPL

print(scipy.optimize.approx_fprime(weights, lambda weights_: fCE(trainX,

```

```

3.5305727448724214e-06
[-1.48039360e-02 -1.48039360e-02 -1.49787697e-02 ...  1.41575085e-04
 8.76898554e-06 -1.38318246e-04]

```

Part(B)

Check_grad weights

The output for these weights is saved in a text file named
"Part_B_check_grad_weights.txt".

```

In [47]: with open("output.txt", "w") as file:
# Gradient check using check_grad
grad_check = scipy.optimize.check_grad(
    lambda weights_: fCE(trainX, trainY, weights_, NUM_HIDDEN_LAYERS,NUM_
    lambda weights_: gradCE(trainX, trainY, weights_, NUM_HIDDEN_LAYERS,
    weights # Initial weights
)

```

```

file.write(f"Gradient check using check_grad: {grad_check}\n\n")

# Gradient check using approx_fprime
approx_grad = scipy.optimize.approx_fprime(weights, lambda weights_: fCE

# Write each value of the approx_fprime output on a new line
file.write("Gradient check using approx_fprime:\n")
for i, value in enumerate(approx_grad):
    file.write(f"Value {i}: {value}\n")

```

```

In [32]: # Next we will perform hyperparameter tuning
         # performed iterations around 162 for different hyperparameters

```

```

In [33]: # Next lets test our model with our best weights and hyperparameters
         test_weights = np.load("best_model_weights_bp.npy")
         test_hyp = np.load("best_model_hyperparameters_bp.npy", allow_pickle=True)

```

```

In [34]: testing_data = np.reshape(np.load("fashion_mnist_test_images.npy"), (-1, 28*
         testing_labels= np.load("fashion_mnist_test_labels.npy")

X_test = testing_data
y_test = testing_labels
X_test = X_test / 255.0 - 0.5

num_input = 784
num_output = 10
hyperparameters = test_hyp.item()
num_hidden_layers = hyperparameters['num_hidden_layers']
num_hidden_units = hyperparameters['num_hidden_units']
num_epochs = hyperparameters['num_epochs']
learning_rate = hyperparameters['learning_rate']
mini_batch = hyperparameters['mini_batch']

num_hidden = num_hidden_layers * [num_hidden_units]

Best_test_accuracy = test(X_test, y_test, test_weights, num_hidden_layers, nu

```

```

In [44]: def show_W0 (W, NUM_HIDDEN_LAYERS, NUM_INPUT, NUM_HIDDEN, NUM_OUTPUT):
         Ws, bs = unpack(W, NUM_HIDDEN_LAYERS, NUM_INPUT, NUM_HIDDEN, NUM_OUTPUT)
         W = Ws[0]
         n = int(NUM_HIDDEN[0] ** 0.5)
         # plt.imshow(np.vstack([
         #     np.hstack([ np.pad(np.reshape(W[idx1*n + idx2,:], [ 28, 28 ]), 2,
         # ]), cmap='gray')

         image_grid = np.vstack([
         np.hstack([
             np.pad(np.reshape(W[idx1 * n + idx2, :], [28, 28]), 2, mode='constant')
             for idx2 in range(n)
         ])
         for idx1 in range(n)])

         # plt.show()

         fig, ax = plt.subplots()

```

```
ax.imshow(image_grid, cmap='gray')

fig.savefig('weight_photo.png', dpi=300, bbox_inches = 'tight')

plt.show()
```

Part (D)

```
In [45]: show_W0(test_weights,num_hidden_layers,num_input,num_hidden,num_output)
```

